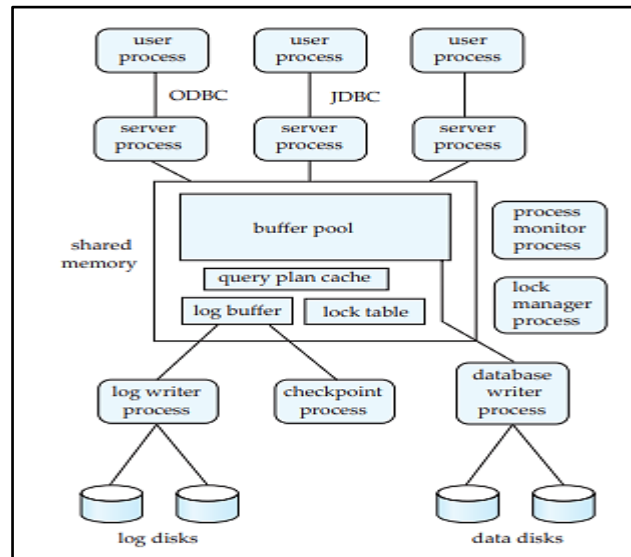


## 1. EXPLAIN TRANSACTION SERVER SYSTEM ARCHITECTURE.

A transaction server system architecture is a computer system that is **designed to process transactions**. A **transaction is a unit of work that is performed on a database**. **Transactions are typically ACID (Atomicity, Consistency, Isolation, and Durability)** compliant, which means that they **must be completed** in their entirety or not at all, and they **must not interfere** with each other.



Transaction server system architectures typically **consist of multiple processes that communicate with each other through shared memory**. The main processes in a transaction server system architecture are:

- **Server processes:** These processes receive user queries and execute them.
- **Lock manager process:** This process manages the locks that are used to ensure that transactions do not interfere with each other.
- **Database writer process:** This process writes modified data from shared memory to disk.
- **Log writer process:** This process writes log records to disk to ensure that the database can be recovered in the event of a failure.
- **Checkpoint process:** This process **periodically takes checkpoints** of the database to reduce the amount of data that needs to be recovered in the event of a failure.
- **Process monitor process:** This process monitors the other processes in the system and takes corrective action if any of them fail.

The shared memory in a transaction server system architecture contains the following data:

- **Buffer pool:** This is a **cache of recently accessed data**.
- **Lock table:** This table contains information about the locks that are held by transactions.
- **Log buffer:** This buffer contains **log records that are waiting to be written to disk**.
- **Cached query plans:** These are **pre-computed plans** for executing common queries.

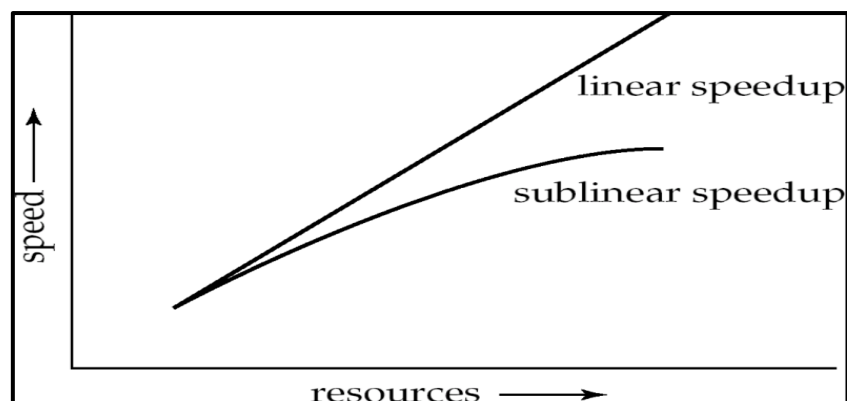
All of the processes in a transaction server system architecture can access the data in shared memory. However, **mutual exclusion mechanisms** are used to ensure that only one process can modify a data at a time.

Transaction server system architectures are **designed to be highly scalable and reliable**. They can **handle large numbers of concurrent transactions** and can **recover from failures quickly**.

## 2. EXPLAIN TERMS.

### 1. Speed up

In the context of a database management system (DBMS), **speed up refers to the improvement in the performance of a database system**. It is typically measured as the **reduction in response time or execution time of queries and transactions**. Speeding up a DBMS involves optimizing various components, such as **query processing, indexing, caching, and parallel processing**, to make the system faster and more efficient. By applying optimizations, the system can handle workloads more swiftly and deliver quicker results to users. Speed up is often quantified as the **ratio of the original performance (before optimizations) to the improved performance (after optimizations)**.



#### Explanation:


The diagram illustrates the concept of speed up in a DBMS.

1. "Original Performance(**Sublinear Speedup**)" represents the initial performance of the database system, which may not be optimal.
2. "Optimization" refers to various **techniques and strategies** applied to enhance the system's performance. These optimizations can include **indexing, query tuning, caching**, and more.
3. As a result of these optimizations, the performance of the database system is improved, leading to "Improved Performance(**Linear Speedup**)."

### 2. Scale up

Scale up involves increasing the capacity and resources of a database system by adding more powerful hardware components to a single server or node. It aims to handle a

greater workload and larger datasets by making the existing server more capable in terms of CPU, memory, and storage.



#### Explanation:

The diagram illustrates the concept of scale up in a DBMS.

1. "Single Server" represents the original configuration of the database system with a single server or node.
2. "Add More Resources" involves upgrading the existing server by increasing CPU, memory, storage, or other hardware resources.
3. This upgrade results in an "Increased Capacity," allowing the system to handle larger workloads and datasets more efficiently.

#### Speedup and scaleup are often sublinear due to several factors:

1. **Start-up costs:** When parallelizing tasks or adding more resources, there can be initial overhead in starting up multiple processes or configuring the new hardware. This start-up cost can dominate the overall computation time, especially when the degree of parallelism is high.
2. **Interference:** In a parallel environment where processes access shared resources, there can be contention and competition among processes. This competition results in waiting time for acquiring shared resources, which reduces the overall efficiency and slows down the speedup or scaleup.
3. **Skew:** High degrees of parallelism can lead to variations in the service times of parallel tasks, known as skew. This variation can be caused by uneven distribution of workload or differences in task complexities. Skew can limit the overall speedup or scaleup as some tasks take longer to complete than others.

Considering these sublinear factors is essential when evaluating the potential improvements in performance or capacity through speedup or scaleup strategies.

In summary, speed up focuses on optimizing various components of a DBMS to improve its performance, while scale up involves increasing the capacity and resources of a database

system by enhancing a single server. Both concepts are crucial for ensuring that a DBMS can meet the demands of growing data and user requirements."

### 3. WHAT ARE COMPLEX OBJECTS? EXPLAIN DIFFERENT COMPLEX TYPES WITH SUITABLE EXAMPLES.

In database management systems (DBMS), *complex objects refer to data types that can encapsulate multiple attributes or components within a single instance.* These complex objects are *used to represent more structured and intricate information* than simple data types like integers or strings. Here are some common complex object types in DBMS:

#### 1. Structured Types:

- **Example:** Consider a "Person" structured type with attributes like name, address, and phone number. Each instance of this type encapsulates a set of related attributes.

```
CREATE TYPE PersonType AS OBJECT (
    name VARCHAR(50),
    address VARCHAR(100),
    phone_number VARCHAR(15)
);
```

#### 2. Arrays:

- **Example:** An "Employee" entity might have an array attribute for a list of project IDs in which the employee is involved.

```
CREATE TYPE EmployeeType AS OBJECT (
    emp_id INT,
    project_ids INT ARRAY
);
```

#### 3. Sets:

- **Example:** A "Course" entity might have a set attribute containing the student IDs of individuals enrolled in that course. The set would represent a collection of student IDs within a single attribute.

```
CREATE TYPE CourseType AS OBJECT (
    course_id INT,
    enrolled_students INT SET
);
```

#### 4. Records:

- **Example:** In a "Library" database, a "Book" entity might have a record attribute containing information such as the book's title, author, and publication date.

```
CREATE TYPE BookInfoType AS OBJECT (
    title VARCHAR(100),
    author VARCHAR(50),
    publication_date DATE
);
```

```
CREATE TYPE BookType AS OBJECT (
    book_info BookInfoType
);
```

## 5. Lists:

- **Example:** A "ToDoList" entity might have a list attribute containing tasks to be completed. The list encapsulates multiple task items within a single attribute.

```
CREATE TYPE ToDoListType AS OBJECT (
    tasks VARCHAR(255) LIST
);
```

These complex object types allow for a more flexible and expressive representation of data in a database, especially when dealing with information that naturally exhibits a structured or nested nature. The specific types available can vary depending on the features and capabilities of the DBMS being used.

## 4. WITH THE USE OF EXAMPLE EXPLAIN TYPE & TABLE INHERITANCE WITH OODBMS.

### Inheritance

It is the process that allows a type or a table to acquire the properties of another type or table. The type or table that inherits the properties is called the subtype or suitable. The type or table whose properties are inherited is called the supertype or super table. Inheritance allows for incremental modification so that a type or table can inherit a general set of properties and add properties that are specific to itself.

You can use inheritance to make modifications only to the extent that the modifications do not alter the inherited supertypes or super tables.

Type and table inheritance are two of the most important features of object-oriented database management systems (OODBMS).

Type inheritance allows you to create new types that inherit the properties and behaviour of existing types. This is similar to class inheritance in object-oriented programming languages.

Table inheritance allows you to create new tables that inherit the columns and constraints of existing tables. This is similar to table inheritance in relational database management systems (RDBMS).

## Example: Table and Type inheritance

```
-- Define the "Person" type
CREATE TYPE Person AS (
    name VARCHAR(100),
    age INT
);

-- Define the "Student" type extending "Person"
CREATE TYPE Student UNDER Person AS (
    major VARCHAR(50),
    minor VARCHAR(50)
);

-- Define the "Teacher" type extending "Person"
CREATE TYPE Teacher UNDER Person AS (
    department VARCHAR(50),
    courses TEXT[] -- An array of course names
);
```

```
-- Table Inheritance Example

-- Create the parent table "Person"
CREATE TABLE Person (
    id SERIAL PRIMARY KEY,
    name VARCHAR(100),
    age INT
);

-- Create the child table "Student" inheriting from "Person"
CREATE TABLE Student (
    student_id SERIAL PRIMARY KEY,
    major VARCHAR(50),
    minor VARCHAR(50),
    FOREIGN KEY (student_id) REFERENCES Person (id)
);

-- Create the child table "Teacher" inheriting from "Person"
CREATE TABLE Teacher (
    teacher_id SERIAL PRIMARY KEY,
    department VARCHAR(50),
    courses TEXT[],
    FOREIGN KEY (teacher_id) REFERENCES Person (id)
);
```

The **Person type is the base type** for the **Student** and **Teacher types**. The Student and Teacher types inherit the name and age properties from the Person type. The Student type also has the major and minor properties, and the Teacher type also has the department and courses properties.

Type and table inheritance can be **used to create complex data models** that reflect the real world. For example, you could use type and table inheritance to model a university database, with different types of students and teachers, and different tables for courses, grades, and other data.

## 5. EXPLAIN CONSTRUCTION OF USER DEFINED DATATYPE.

User-Defined Types (UDT) in a database allow you to create custom data types tailored to your specific needs. Here's a brief explanation:

### 1. What is UDT?

- **Definition:** UDT refers to User-Defined Types, which are custom data types created by users based on their specific data structure requirements.
- **Purpose:** UDT allows you to define a set of related attributes as a single logical unit, enhancing data organization and readability.

### 2. Why is UDT Required?

- **Customization:** UDT enables the creation of data types tailored to specific business requirements or data structures that are not adequately represented by standard data types.
- **Encapsulation:** UDT allows you to encapsulate related attributes into a single type, promoting modularity and simplifying the database schema.

### 3. What UDT Does:

- **Structured Data:** UDT helps in representing structured or composite data where a group of related attributes belongs together.
- **Improved Readability:** UDT enhances the readability of your database schema by grouping related attributes into meaningful units, making it easier to understand and maintain.
- **Data Integrity:** By defining a clear structure for data, UDT contributes to maintaining data integrity and consistency.

## 6. EXPLAIN THE CONCEPT OF REFERENCING IN OODBMS. [UNIT I]

In Object-Oriented Database Management Systems (OODBMS), referencing allows objects to establish relationships with other objects. This enables navigation between related objects and facilitates efficient data retrieval. Key points about referencing in OODBMS are:

- Referencing establishes relationships between objects.
- Objects can refer to other objects using unique identifiers or object pointers.
- Referencing enables navigation and access to related objects directly.
- It helps maintain the integrity of the object and ensures data consistency.
- Object identity and uniqueness are established through referencing.
- Referencing supports object persistence, allowing objects to be stored and retrieved from the database.

Let's consider an example of referencing in an Object-Oriented Database Management System (OODBMS) using a "Student" and "Course" scenario:

**class Student:**

**attributes:**

student\_id

name

age

courses # List to represent the many-to-many relationship with Course objects

**class Course:**

**attributes:**

course\_id

title

instructor

**# Creating instances for students**

john = new Student(student\_id=1, name='John', age=20)

jane = new Student(student\_id=2, name='Jane', age=22)

**# Creating instances for courses**

```

mathematics = new Course(course_id=1, title='Mathematics', instructor='Professor Smith')
english_literature = new Course(course_id=2, title='English Literature', instructor='Professor Johnson')

```

### # Establishing references

```

john.courses.add(mathematics)
john.courses.add(english_literature)
jane.courses.add(english_literature)

```

In this example, referencing allows us to establish relationships between students and courses. Students can enrol in multiple courses, and each course can have multiple students.

## 7. COMPARE OODBMS WITH ORDBMS. [UNIT I]

### OODBMS

It is an extension of an object-oriented programming language that includes DBMS functions such as ***persistent objects, integrity constraints, failure recovery, transaction management, and query processing.***

Some examples of OODBMS are ***ObjectStore, Objectivity/DB, GemStone, db4o, Giga Base, and Zope object database.***

### ORDBMS

It is a ***relational database system that incorporate object-oriented characteristics.*** Database ***schemas*** and the ***query language*** natively ***support objects, classes, and inheritance.***

***Oracle, DB2, Informix, PostgreSQL*** etc. are some of the ORDBMS.

Feature	OODBMS	ORDBMS
<b>Data model</b>	Object-oriented	Relational
<b>Data representation</b>	Objects, Classes	Tables( Attributes, Records )
<b>Relationships</b>	Object references, inheritance	Foreign keys
<b>Queries</b>	Object query language (OQL)	Structured query language (SQL)
<b>Scalability</b>	Horizontal scalability is easier	Vertical scalability is easier
<b>ACID compliance</b>	Yes	Yes
<b>Maturity</b>	Less mature	More mature
<b>Use cases</b>	<b><i>Complex data structures, multimedia data</i></b>	<b><i>Structured data, transactional data</i></b>
<b>Example</b>	Objectivity/DB, ObjectBase, GigaBase, Zope Object Store etc.	Oracle, PostgreSQL, DB2 etc.

## 8. WITH BLOCK DIAGRAM EXPLAIN DISTRIBUTED DATABASE SYSTEM ARCHITECTURE.



In a distributed system architecture, the **database is distributed across multiple computers, referred to as sites or nodes**, which communicate with each other through various communication media like **high-speed private networks or the Internet**. Unlike shared-memory or shared-disk architectures, **computers in a distributed system do not share main memory or disks**. The structure of a distributed system involves various sites with varying sizes and functions, ranging from **workstations to mainframe systems**.

#### **Key Characteristics:**

1. **Geographical Separation:** Distributed databases are typically geographically separated, requiring communication over networks.
2. **Separate Administration:** Each distributed database site is separately administered, allowing a degree of local control and autonomy.
3. **Local and Global Transactions:** Distributed databases differentiate between local and global transactions. Local transactions access data only from the site where the transaction was initiated, while global transactions access data in different sites.

#### **Reasons for Building Distributed Database Systems:**

1. **Sharing Data:** Enables users at one site to access data residing at other sites, **facilitating collaboration and eliminating the need for external mechanisms**.
2. **Autonomy:** Autonomy refers to the **degree of control or independence** that each local site has over its own database. Local sites give up some of this autonomy when it comes to **making changes to the database**.
3. **Availability:** Offers increased availability by allowing the system to continue operating even if one site fails, particularly when data items are replicated across multiple sites.

#### **Example of a Distributed Database:**

Consider a **banking system** with four branches in different cities, each having its own computer and database. One site maintains information about all the branches. Transactions can be local (e.g., adding money to an account at the branch where the transaction was initiated) or global (e.g., transferring money between accounts in different branches).




Figure 17.9 A distributed system.

### Implementation Issues:

1. **Atomicity of Transactions:** Atomicity ensures that the transaction is treated as a single, and all involved branches either commit the transaction or abort it consistently, maintaining data integrity across the distributed database system. The two-phase commit protocol (2PC) is commonly used.
2. **Concurrency Control:** Concurrency control refers to managing and coordinating multiple transactions that are executing simultaneously in a distributed database system. It ensures that transactions don't interfere with each other and maintain data consistency.
3. **Failures and Recovery:** Distributed systems must handle failures, including computer failures and communication link failures. Recovery mechanisms are crucial to restore normal operation.

### Challenges:

1. **Increased Complexity:** Building and managing distributed database systems involve added complexity, including higher software development costs, greater potential for bugs, and increased processing overhead due to intersite coordination and communication.

In summary, distributed system architecture aims to provide a collaborative environment for data sharing while addressing challenges related to geographical separation, separate administration, and increased complexity.

### 9. ELABORATE OBJECT REFERENCING IN ORDBMS. WRITE A SCHEMA DEFINITION CORRESPONDING TO FOLLOWING RELATIONAL SCHEMA USING REFERENCES: [UNIT 1]

**Emp**(Person nm, street, city), **Works**(Person nm, Comp\_nm, salary), **Company**(Comp\_nm, city), **Manages**(Person\_nm, manager nm)

```

CREATE TABLE Emp (
  Person_nm VARCHAR(255) NOT NULL,
  street VARCHAR(255),
  city VARCHAR(255),
  PRIMARY KEY (Person_nm)
);

CREATE TABLE Works (
  Person_nm VARCHAR(255) NOT NULL,
  Comp_nm VARCHAR(255) NOT NULL,
  salary INT,
  PRIMARY KEY (Person_nm, Comp_nm),
  FOREIGN KEY (Person_nm) REFERENCES Emp(Person_nm)
);

CREATE TABLE Company (
  Comp_nm VARCHAR(255) NOT NULL,
  city VARCHAR(255),
  PRIMARY KEY (Comp_nm)
);

CREATE TABLE Manages (
  Person_nm VARCHAR(255) NOT NULL,
  manager_nm VARCHAR(255) NOT NULL,
  PRIMARY KEY (Person_nm),
  FOREIGN KEY (Person_nm) REFERENCES Emp(Person_nm),
  FOREIGN KEY (manager_nm) REFERENCES Emp(Person_nm)
);

```

## 1. WHAT ARE THE DIFFERENT TYPES OF DISTRIBUTED DB SYSTEMS? DIFFERENTIATE BETWEEN HOMOGENEOUS AND HETEROGENEOUS DATABASES. UNIT [2]

### Types of Distributed Database Systems:

#### 1. Homogeneous Distributed Database System:

- Characteristics:
  - All sites have identical database management system (DBMS) software.
  - Sites are aware of each other and agree to cooperate in processing user requests.
  - Supports Autonomy.

- Cooperation is necessary for exchanging information about transactions to enable processing across multiple sites.

- **Advantages:**

- Simplicity in terms of software consistency.
- Easier coordination and communication between sites.

## 2. Heterogeneous Distributed Database System:

- **Characteristics:**

- Different sites may use different schemas and DBMS software.
- Sites may not be aware of each other and may provide limited facilities for cooperation in transaction processing.
- Challenges arise due to differences in schemas, data models, and software.

- **Advantages:**

- Allows for diverse data models and local autonomy.
- Flexibility in accommodating existing databases without requiring significant changes.

### Differences between Homogeneous and Heterogeneous Databases:

Characteristic	Homogeneous Distributed Database	Heterogeneous Distributed Database
<b>Software Consistency</b>	All sites have identical DBMS software.	Different sites may use different schemas and DBMS software.
<b>Awareness and Cooperation</b>	Sites are aware of each other and cooperate in processing user requests.	Sites may not be aware of each other and provide limited cooperation.
<b>Autonomy</b>	Local sites surrender some autonomy in terms of schema changes or DBMS software modifications.	Local sites retain a high degree of autonomy over their schemas and transactions.
<b>Data Model</b>	Conforms to a common data model.	May involve integration of diverse data models into a common one.
<b>Flexibility</b>	Limited flexibility due to the need for consistency.	Allows for flexibility in accommodating diverse data models and local requirements.


In summary, homogeneous distributed databases aim for consistency through uniformity, while heterogeneous distributed databases accommodate diversity in terms of data models, schemas, and local autonomy. Each type has its advantages and challenges, and the choice depends on factors such as system requirements, existing databases, and organizational considerations.

## 2. EXPLAIN THE ARCHITECTURE OF DISTRIBUTED TRANSACTION SYSTEM. [UNIT 2]

In a distributed transaction system, the architecture consists of multiple sites, each with its own local transaction manager and transaction coordinator subsystems. The purpose of this architecture is to ensure the ACID (Atomicity, Consistency, Isolation, Durability) properties of both local and global transactions.

### 1. Local Transaction Manager:

Each site in the distributed system has its own local transaction manager responsible for managing the execution of transactions that access data stored locally. The local transaction manager ensures the ACID properties of transactions executing at that site. It performs tasks such as maintaining a log for recovery purposes and participating in concurrency-control schemes to coordinate the concurrent execution of transactions at the site.



### 2. Transaction Coordinator:

The transaction coordinator subsystem is responsible for coordinating the execution of all transactions initiated at a particular site. It oversees global transactions that involve accessing and updating data across multiple sites. The transaction coordinator performs the following tasks:

- **Starting the execution of the transaction:** The coordinator initiates the execution of a global transaction.
- **Breaking the transaction into sub transactions:** The coordinator divides the transaction into sub transactions and distributes them to the appropriate sites for execution.
- **Coordinating the termination of the transaction:** The coordinator ensures that the global transaction is either committed at all participating sites or aborted at all sites.

The overall system architecture includes multiple transaction managers and transaction coordinators, as illustrated in Figure. Each site has its own local transaction manager (TM1, TM2, ...,

TMn) and a transaction coordinator (TC1, TC2, ..., TCn) responsible for coordinating transactions at that site.

### 3. System Failure Modes:

Distributed systems face various types of failures beyond those encountered in centralized systems. These failure modes include:

- **Failure of a site:** A site may experience software errors, hardware failures, or disk crashes, leading to unavailability of data or services.
- **Loss of messages:** Messages exchanged between sites may be lost or corrupted during transmission. Transmission-control protocols like TCP/IP are used to handle such errors.
- **Failure of a communication link:** If a communication link between two sites fails, messages must be rerouted or alternate routes found to ensure communication between sites.

To handle these failure modes, distributed systems employ fault-tolerance mechanisms, recovery protocols, and communication protocols to maintain data integrity, ensure reliable message delivery.

## 3. ELABORATE IN DETAIL DATA REPLICATION AND DATA FRAGMENTATION IN DISTRIBUTED DATABASES. [UNIT 2]

Data replication and data fragmentation are two important techniques used in distributed databases to improve performance, scalability, and availability.

Data replication involves storing multiple copies of the same data at different sites in the distributed database system. This can improve read performance by reducing the distance that data needs to be travelled to be accessed. It can also improve availability by ensuring that data is still available even if one site in the system fails.

**There are two main types of data replication:**

- **Full replication:** All data is replicated at all sites in the system. This is the simplest type of replication, but it can be expensive and complex to implement.
- **Partial replication:** Only a subset of the data is replicated at each site. This can be more efficient than full replication, but it can be more complex to manage and can lead to consistency problems if the data is not replicated correctly.

### Data Fragmentation:

Fragmentation is a process of dividing the whole or full database into various sub tables or sub relations so that data can be stored in different systems. The small pieces or sub relations or sub tables are called fragments. These fragments are called logical data units and are stored at various sites. It must be made sure that the fragments are such that they can be used to reconstruct the original relation (i.e., there isn't any loss of data).

**Fragmentation can be of three types:**

- **Horizontal fragmentation:** Horizontal fragmentation refers to the process of dividing a table horizontally by assigning each row (or a group of rows) of relation to one or more fragments.
- **Vertical fragmentation:** Vertical fragmentation involves creating tables with fewer columns and using additional queries to save disk space.
- **Mixed or Hybrid fragmentation:** Mixed fragmentation is a combination of horizontal and vertical fragmentation.

Both data replication and fragmentation are crucial for the efficient functioning of distributed databases. They help in improving data availability, enhancing system performance, and ensuring data consistency across different sites.

#### 4. EXPLAIN THE 2-PHASE COMMIT PROTOCOL FOR DISTRIBUTED DATABASE RECOVERY. [UNIT 2]

##### Two-Phase Commit Protocol for Distributed Database Recovery:

The Two-Phase Commit Protocol (2PC) is a mechanism used in distributed databases to ensure the consistency of transactions across multiple sites. It operates in three main phases: the commit protocol, handling of failures, and recovery with concurrency control.

##### 1. The Commit Protocol:

- **Phase 1: Prepare**
  - After a transaction T completes its execution at a site  $S_i$ , the coordinator  $C_i$  initiates Phase 1 of the 2PC.
  - $C_i$  adds a record  $\langle \text{prepare } T \rangle$  to the log and forces the log onto stable storage.
  - $C_i$  sends a prepare T message to all sites where T executed.
  - Each participating site responds:
    - If willing to commit, it adds  $\langle \text{ready } T \rangle$  to the log and forces the log onto stable storage. The site replies with ready T to  $C_i$ .
    - If not willing to commit, it adds  $\langle \text{no } T \rangle$  to the log, sends an abort T message to  $C_i$ , and does not proceed to Phase 2.
- **Phase 2: Commit or Abort**
  - When  $C_i$  receives responses or a specified time elapse, it determines whether to commit or abort T.
  - If all sites responded with ready T,  $C_i$  adds  $\langle \text{commit } T \rangle$  to the log; otherwise, it adds  $\langle \text{abort } T \rangle$ .
  - The decision is forced onto stable storage.
  - $C_i$  sends either commit T or abort T messages to all participating sites.

- Sites record the message in the log, and T enters the ready state at each site.
- **Acknowledgment (optional):**
  - Some implementations include an acknowledge T message from sites to Ci at the end of Phase 2.
  - When Ci receives acknowledgments from all sites, it adds <complete T> to the log.

## 2. Handling of Failures:

- **Failure of a Participating Site:**
  - Ci assumes abort if a site fails before responding with ready T.
  - After recovery, the site determines T's fate based on log records (commit, abort, or ready).
- **Failure of the Coordinator:**
  - Active sites decide T's fate based on their logs.
  - If the coordinator failed, and no <ready T> record is found, active sites must wait for recovery.
  - The blocking problem arises, causing potential delays in deciding T's fate.
- **Network Partition:**
  - The protocol accommodates network partitions without affecting normal operation.

## 3. Recovery and Concurrency Control:

- Recovery involves determining the status of in-doubt transactions (missing <commit T> or <abort T> records).
- Recovery algorithms may use lock information in the log, allowing faster restart recovery.
- Instead of <ready T>, the algorithm writes <ready T, L> records, including a list of write locks held by T.
- Lock reacquisition is performed during recovery, and normal transaction processing can resume concurrently with commit or rollback of in-doubt transactions.

The 2PC protocol ensures distributed transaction consistency, handles failures, and supports recovery with concurrency control, although the blocking problem during coordinator failure remains a challenge.

## 5. GIVE DETAILS OF LOCKING PROTOCOL IN DISTRIBUTED DATABASES/ ELABORATE IN DETAIL IMPLEMENTATION OF LOCKING PROTOCOL IN DISTRIBUTED DATABASES. [UNIT 2]

In distributed databases, a locking protocol is used to ensure the consistency and concurrency control of data when multiple transactions are accessing and modifying the data concurrently. The



provided information describes different locking protocols that can be applied in a distributed environment. Here are the details of the locking protocols mentioned:

1. **Single Lock-Manager Approach:**

In this approach, there is a single lock manager that resides in a chosen site. All lock and unlock requests are made at that site. When a transaction needs to lock a data item, it sends a lock request to the lock manager. The lock manager determines if the lock can be granted immediately or if it needs to be delayed. The transaction can read the data item from any replica site, but all replica sites must be involved in writing. Advantages include simple implementation and deadlock handling, but it suffers from a bottleneck at the chosen site and vulnerability if that site fails.

2. **Distributed Lock Manager:**

This approach distributes the lock-manager function across multiple sites. Each site maintains a local lock manager responsible for administering lock and unlock requests for the data items stored in that site. When a transaction wants to lock a data item, it sends a lock request to the lock manager at the respective site. The lock manager determines if the lock can be granted immediately and sends a response back to the transaction. This approach reduces the bottleneck but introduces complexity in deadlock handling, as intersite deadlocks may occur.

1. **Primary Copy:**

In this approach, one replica of a data item is designated as the primary copy, and it resides in a single site. When a transaction needs to lock the data item, it requests a lock at the primary site. This approach simplifies concurrency control, but if the primary site fails, the data item becomes inaccessible, even if other replica sites are available.

2. **Majority Protocol:**

The majority protocol requires a lock-request message to be sent to more than half of the sites where a data item is replicated. Each lock manager determines if it can grant the lock request immediately. The transaction operates on the data item only after obtaining a lock on a majority of the replicas. This protocol can deal with site failures and offers decentralized control but requires more complex implementation and deadlock handling.

3. **Quorum Consensus Protocol:**

The quorum consensus protocol extends the majority protocol by assigning each site a weight. Read and write operations are assigned read quorum and write quorum, respectively, based on the weights of the sites where the data item resides. The total weight of the locked replicas must satisfy certain conditions for read and write operations to proceed. This protocol allows flexibility in setting read and write quorums and can simulate other protocols by appropriately defining weights and quorums.

These locking protocols provide various trade-offs in terms of simplicity, concurrency control, scalability, and fault tolerance in a distributed database environment. The choice of protocol depends on the specific requirements and characteristics of the system.

## 6. EXPLAIN DEADLOCK HANDLING IN DISTRIBUTED DATABASES. [UNIT 2]

In a distributed database, deadlock handling refers to the techniques and algorithms used to detect and resolve deadlocks that can occur when multiple transactions compete for resources across different sites. Deadlocks occur when two or more transactions are waiting for resources held by each other, resulting in a circular dependency that prevents any transaction from progressing.

There are two main approaches to deadlock handling in distributed databases: deadlock prevention and deadlock detection. Let's explore each of these approaches in more detail:

### 1. **Deadlock Prevention:**

Deadlock prevention aims to eliminate the conditions that can lead to a deadlock by carefully managing resource allocation. This approach requires coordination among the sites to ensure that transactions acquire resources in a way that avoids circular dependencies. Some common techniques for deadlock prevention include:

- **Resource Ordering:** Transactions are required to request resources in a specified order to prevent circular wait conditions. This technique establishes a global order for acquiring resources based on predefined rules, such as assigning a unique identifier to each resource and enforcing transactions to request resources in ascending order of their identifiers.
- **Wait-Die and Wound-Wait:** These techniques involve assigning priorities or timestamps to transactions. In the Wait-Die approach, a younger transaction waits for an older transaction to release a resource, while in the Wound-Wait approach, an older transaction forcibly aborts a younger transaction to free up resources. These techniques prevent circular waits by enforcing a strict order of resource allocation.
- **Two-Phase Locking:** Two-Phase Locking (2PL) is a concurrency control technique where transactions acquire and release locks on resources in two phases: the growing phase and the shrinking phase. By adhering to the strict rules of 2PL, deadlocks can be prevented.


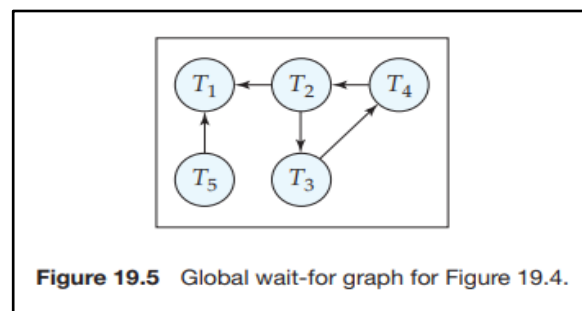

Deadlock prevention techniques can be effective in avoiding deadlocks, but they may introduce additional waiting and rollback operations, and may require more sites to be involved in transaction execution than necessary.

### 2. **Deadlock Detection:**

Deadlock detection focuses on identifying the existence of deadlocks when they occur. Rather than preventing deadlocks, this approach allows them to happen and then uses algorithms to detect and resolve them. In a distributed system, maintaining a wait-for graph is a common technique for deadlock detection. The wait-for graph represents the dependencies between transactions and resources.

In a distributed environment, each site maintains a local wait-for graph that includes both local and non-local transactions requesting or holding resources at that site. Whenever a transaction requests a resource from another site, a request message is sent, and if the requested resource is held by another transaction, an edge is added to the local wait-for graph of the site holding the resource.

To detect deadlocks, a centralized deadlock-detection coordinator periodically collects and combines the local wait-for graphs from all sites to construct a global wait-for graph. If the global wait-for graph contains a cycle, deadlock is detected. Once a deadlock is detected, a victim transaction is selected and rolled back, breaking the circular dependency.



However, the global wait-for graph can sometimes contain false cycles or may lead to unnecessary rollbacks. False cycles can occur due to communication delays or timing issues, leading to incorrect deadlock detection. Additionally, a victim transaction may be rolled back even if it was not directly involved in the deadlock due to unrelated reasons, causing unnecessary rollback operations.

Overall, deadlock handling in distributed databases involves a combination of prevention techniques to reduce the occurrence of deadlocks and detection techniques to identify and resolve deadlocks when they do occur. The choice of approach depends on the specific requirements and trade-offs of the distributed database system.

## 7. WITH A SUITABLE EXAMPLE ELABORATE DISTRIBUTED LOCK MANAGEMENT TO CONTROL CONCURRENCY. [UNIT 2]

Distributed lock management is an essential component of concurrency control in distributed databases. It ensures that multiple transactions accessing shared resources across different sites do so in a coordinated and controlled manner, preventing conflicts and maintaining data consistency. Let's illustrate distributed lock management with an example:

Consider a distributed database system with two sites, Site A and Site B. Each site has a local lock manager responsible for administering lock and unlock requests for the data items stored at that site.

1. Transaction T1 at Site A wants to access data item D1. It sends a lock request to the lock manager at Site A.
2. The lock manager at Site A checks if D1 is already locked.
3. If D1 is not locked, the lock manager grants the lock to T1.
4. Now, suppose Transaction T2 at Site B also wants to access D1. It sends a lock request to the lock manager at Site B.
5. The lock manager at Site B communicates with the lock manager at Site A and finds out that D1 is already locked by T1.
6. The lock manager at Site B then places T2 in a waiting state until T1 releases the lock on D1.
7. Once T1 is done with D1, it sends an unlock request to the lock manager at Site A.
8. The lock manager at Site A releases the lock and communicates this to the lock manager at Site B.
9. The lock manager at Site B can now grant the lock on D1 to T2.

This is a simplified example, but it illustrates the basic workings of distributed lock management. In reality, the process can be much more complex, involving multiple lock managers, different types of locks (shared, exclusive), and sophisticated algorithms to handle deadlocks.

## **8. EXPLAIN CONCURRENCY CONTROL SCHEMES IN DISTRIBUTED DATABASES. [UNIT 2]**

Concurrency control in distributed databases is crucial to ensure that transactions executed concurrently maintain consistency and isolation. Several concurrency control schemes are employed to manage the access and modification of shared data across multiple sites. Here are some common concurrency control schemes in distributed databases:

### **1. Locking:**

- **Centralized Lock Manager:**

- A central server manages locks for all distributed transactions.
- Sites request locks from the central manager.
- Ensures global consistency but may introduce bottlenecks.

- **Distributed Lock Manager:**

- Each site manages its locks, and coordination is required for global consistency.
- Allows for more parallelism but requires additional coordination.

## 2. Timestamp-based Protocols:

- **Centralized Timestamps:**

- A central authority assigns timestamps to transactions.
- Timestamps used to determine the order of execution and resolve conflicts.
- Centralized control may become a bottleneck.

- **Distributed Timestamps:**

- Each site generates its timestamps.
- Coordination is required to maintain a consistent global order of transactions.

## 3. Optimistic Concurrency Control:

- **Timestamp Ordering:**

- Each transaction is assigned a timestamp.
- Transactions can proceed without obtaining locks based on their timestamps.
- At commit time, conflicts are detected, and the commit is aborted if necessary.

- **Validation-based:**

- Transactions are executed without locks or coordination.
- Validation phase checks for conflicts before committing.
- If conflicts are detected, the transaction is rolled back.

## 4. Multi-version Concurrency Control (MVCC):

- Each transaction sees a snapshot of the database as of the transaction's start time.
- Multiple versions of data may exist, allowing for concurrent reads and writes without conflicts.
- Ensures isolation without locking, improving parallelism.

## 5. Two-Phase Locking (2PL):

- Applies the principles of 2PL at a global level in a distributed environment.
- Transactions acquire and release locks in two phases (growing phase and shrinking phase) to maintain consistency.

## 6. Quorum-based Protocols:

- Transactions need to obtain a quorum of votes from a subset of sites.
- Quorum can be configured to ensure consistency and fault tolerance.
- Can be used in systems with replication.

## 7. Consistency Models:

- **Eventual Consistency:**
  - Allows for temporary inconsistencies, resolving over time.
  - Suitable for systems where eventual convergence is acceptable.
- **Causal Consistency:**
  - Enforces causality between related transactions.
  - Provides a balance between consistency and performance.

These concurrency control schemes aim to balance the trade-offs between consistency, isolation, and performance in distributed database systems

## 9. EXPLAIN THE FOLLOWING TERMS OF QUERY PROCESSING WITH SUITABLE EXAMPLE.

### 1. SIMPLE JOIN

### 2. SEMI JOIN [UNIT 2]

#### 1. Simple Join:

A simple join refers to the process of combining two or more relations based on a common attribute or condition. It is a fundamental operation in relational databases for retrieving data from multiple tables. The join operation matches the rows from different tables that satisfy the join condition and produces a result set that combines the selected rows.

#### Example:

Consider two relations, "Employees" and "Departments." The "Employees" relation contains information about employees, including their employee ID, name, and department ID. The "Departments" relation stores details about different departments, including the department ID and department name.

To retrieve the employee names along with their corresponding department names, a simple join operation can be performed on the "Employees" and "Departments" relations, using the common attribute "department ID." The SQL query for this would be:

```
SELECT Employees.Name, Departments.DepartmentName
FROM Employees
JOIN Departments ON Employees.DepartmentID = Departments.DepartmentID;
```

The result of the join operation would be a table that includes employee names and their respective department names, based on the matching department IDs.

#### 2. Semi Join:

A semi join is a specialized join operation that aims to reduce the amount of data transmitted over the network when performing a join operation between two relations. It

selectively transfers only the necessary data from one site to another, based on a common attribute, to optimize the overall query processing performance.

The semi join operator ( $\Join$ ) in relational algebra returns the tuples from one relation that have matching values in the other relation, without duplicating the tuples from the second relation.

### Example:

Let's say we have two relations, "Orders" and "Customers." The "Orders" relation contains information about individual orders, including the order ID, customer ID, and order details. The "Customers" relation stores customer information, including the customer ID and customer name.

Suppose we want to retrieve a list of customers who have placed orders. Instead of transferring the entire "Customers" relation to the site where the join operation is performed, a semi join strategy can be employed.

**The SQL query for this would be:**

```
SELECT *
FROM Customers
WHERE CustomerID IN (SELECT CustomerID FROM Orders);
```

In this query, the subquery `(SELECT CustomerID FROM Orders)` performs the semi join operation. It selects only the unique customer IDs from the "Orders" relation. The outer query then retrieves all the customer details from the "Customers" relation, where the customer ID is present in the result of the subquery.

By using a semi join, we avoid transferring unnecessary data from the "Customers" relation to the site where the join operation is performed, resulting in improved query performance, especially when dealing with large datasets and distributed databases.

## 10. EXPLAIN DISTRIBUTED QUERY PROCESSING. [UNIT 2]

Distributed Query Processing is the process of executing a database query that involves data stored across multiple distributed databases. The goal is to **retrieve and integrate relevant data from different locations to provide a unified result.** Let's walk through a simple example to illustrate distributed query processing:

**Scenario:** Consider two geographically distributed databases, Database A and Database B, each storing information about employees.

**Query:** Retrieve the **names and salaries** of all employees who work in the "IT" department.

### Distributed Query Processing Steps:

#### 1. Query Decomposition:

- The query is decomposed into subqueries that can be executed locally at each database. For our example, it might involve extracting employee names and salaries from the "IT" department.

## 2. Subquery Execution:

- Subqueries are sent to the respective databases for execution. Database A executes the subquery to retrieve names and salaries of IT department employees stored locally, and Database B does the same.

## 3. Data Transfer:

- The results of the subqueries are transferred to a central location, often the querying site or another designated node. This transfer can be asynchronous or synchronous, depending on the system.

## 4. Data Integration:

- The results from Database A and Database B are integrated at the central location. This may involve merging or aggregating the data to provide a unified response to the original query.

## 5. Result Presentation:

- The final integrated result is presented to the user or application that initiated the query.

### Original Query

```
SELECT EmployeeName, Salary
FROM Employees
WHERE Department = 'IT';
```

Let's illustrate the process with SQL-like queries for each step:

- **Query Decomposition:**

#### **Subquery for Database A**

```
SELECT EmployeeName, Salary
FROM DatabaseA.Employees
WHERE Department = 'IT';
```

#### **Subquery for Database B**

```
SELECT EmployeeName, Salary
```



```
FROM DatabaseB.Employees
```

```
WHERE Department = 'IT';
```

- **Subquery Execution:**

- The subqueries are executed at their respective databases.

- **Data Transfer:**

- The results are transferred to the central location.

- **Data Integration:**

**Merging results from Database A and Database B**

```
SELECT * FROM ResultA UNION ALL SELECT * FROM ResultB;
```

- **Result Presentation:**

- The final result is presented to the user or application.

## 1. WHAT IS I/O PARALLELISM? EXPLAIN IN DETAIL.[UNIT 3]

I/O parallelism refers to the technique of enhancing the efficiency of retrieving relations from disk by distributing and managing data across multiple disks in a parallel database system. This is achieved through various data partitioning strategies, commonly used to improve the performance of reading and writing data.

### 1. Basic Partitioning Techniques:

- Round-robin Partitioning:
  - Tuples are distributed across disks in a cyclic order.
  - Ensures an even distribution of tuples across disks.
- Hash Partitioning:
  - Uses a hash function on specified attributes to assign tuples to disks.
  - Efficient for point queries based on partitioning attributes.
- Range Partitioning:

- Divides tuples based on contiguous attribute-value ranges.
- Suited for point and range queries on the partitioning attribute.

## 2. Comparison of Partitioning Techniques:

- Different partitioning techniques support various types of data access more efficiently:
  - Round-robin is ideal for sequential scans but complicates point and range queries.
  - Hash partitioning excels at point queries on partitioning attributes and sequential scans.
  - Range partitioning is well-suited for point and range queries on the partitioning attribute.

## 3. Handling Skew:

- Skew can occur due to uneven distribution of tuples, affecting performance.
- Attribute-value skew results from tuples sharing the same partitioning attribute value.
- Partition skew occurs when there's a load imbalance, even without attribute skew.
- Techniques like sorting or using histograms can address skew issues in range partitioning.
- Virtual processors can distribute workload more evenly, mitigating the impact of skew.

## 4. Impact on Performance:

- I/O parallelism significantly enhances transfer rates for reading or writing entire relations.
- Efficient for accessing data sequentially but impacts different partitioning techniques differently for point and range queries.
- Careful consideration of partitioning technique is crucial based on the nature of queries and operations.


In summary, I/O parallelism optimizes data retrieval from disk in parallel database systems by intelligently partitioning and managing data across multiple disks, considering the nature of queries and potential skew issues.

## 2. EXPLAIN DIFFERENT ARCHITECTURES OF PARALLEL DATABASE SYSTEM AND COMPARE AMONG THEM. [UNIT 3]

### Parallel Database Architectures


1. **Shared memory** -- processors share a common memory
2. **Shared disk** -- processors share a common disk
3. **Shared nothing** -- processors share neither a common memory nor common disk
4. **Hierarchical** -- hybrid of the above architectures

#### 1. Shared Memory




- Processors and disks have **access to a common memory**, typically **via a bus** or through an **interconnection network**.
- This architecture provides **efficient communication in between processors** because **data of shared memory can be directly accessed by each processor**.
- Downside** – **architecture is not scalable beyond 32 or 64 processors**.
- Widely used for **lower degrees of parallelism (4 to 8)**.

## 2. Shared Disk



- All processors can directly access all disks via an interconnection network, but the processors have private memories.
- The **memory bus is not a bottleneck**.
- Architecture **provides a degree of fault-tolerance**: If a processor fails, the other processors can take over its tasks.
- Examples**: **IBM Sysplex** and **DEC clusters** running Rdb were early commercial users.
- Downside**: **Bottleneck now occurs at interconnection to the disk** subsystem.
- Shared-disk systems can **scale to a somewhat larger number of processors**, but **communication between processors is slower**.

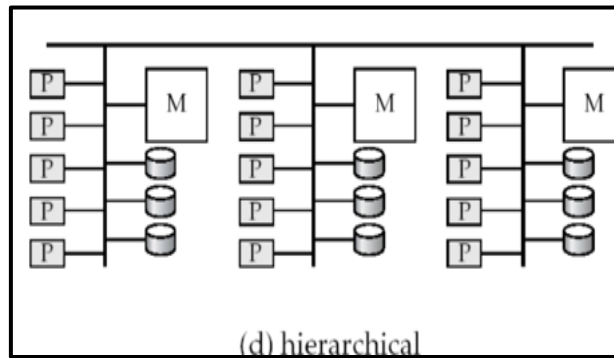
## 3. Shared Nothing



- Each node consists of a **processor, memory, and one or more disks**. Processors at one node communicate with another processor at another node using an interconnection network. A **node functions as the server** for the data on the disk the node owns.
- Examples: **Teradata, Tandem, Oracle-n CUBE**
- Data accessed from local disks do not pass-through interconnection network**.

- Shared-nothing Architecture can be **scaled up to thousands of processors** without interference.
- **Main drawback: cost of communication and non-local disk access;** sending data **involves software interaction** at both ends

#### 4. Hierarchical



- Hierarchical parallel database architectures are well-suited for a wide range of applications, including **data warehousing, business intelligence, scientific computing,** and **enterprise resource planning (ERP).**
- Hierarchical parallel database architecture is a **hybrid architecture that combines the characteristics of shared memory, shared disk, and shared nothing architectures.**
- **At the top level,** a hierarchical parallel database system is a **shared nothing system.** This means that each node in the system has its own CPU, memory, and disk. **The nodes are interconnected by a high-speed network.**
- **At the lower level,** each node in a hierarchical parallel database system can be a **shared memory or shared disk system.** This means that the processors on each node can share memory and/or disks.
- Hierarchical parallel database architectures are designed to **achieve high scalability and performance.**

### 3. TYPES OF PARTITIONING IN PARALLEL DATABASE. [UNIT 3]

Partitioning in parallel databases is the **process of dividing the data into multiple partitions,** which can then be **processed in parallel by processing units** of that server.

This can **improve the performance of query processing** of the database system.

There are two main types of partitioning in parallel databases:

1. horizontal partitioning

2. vertical partitioning

#### 1. HORIZONTAL PARTITIONING

Horizontal partitioning divides the data into multiple tables, where **each table contain subset of the rows.** This is the most common type of partitioning in parallel

databases, as *it allows for the data to be processed in parallel across multiple processing units.*

**There are several different ways to perform horizontal partitioning, including:**

**1. Round Robin:** It is a type of horizontal partitioning in which the *data is evenly distributed across multiple partitions in a round robin fashion or in circular fashion.*

Round robin partitioning is a *simple and efficient way to distribute data evenly.* It is particularly well-suited for workloads where the data is not evenly distributed.

Here's how round-robin partitioning works with an example:

Suppose you have a database table containing customer data with a unique customer ID as the primary key. You want to partition this table across four processing nodes/units in a parallel database system.

**Data Partitioning:**

- The customer data is divided into four partitions (P1, P2, P3, and P4).
- To achieve round-robin partitioning, you assign each customer ID to a partition in a circular fashion, starting with the first partition and cycling through all partitions until all customer IDs are assigned.

**Example:**

- Customer ID 1 goes to Partition P1.
- Customer ID 2 goes to Partition P2.
- Customer ID 3 goes to Partition P3.
- Customer ID 4 goes to Partition P4.
- Customer ID 5 goes back to Partition P1.
- This pattern continues until all customer IDs are assigned to partitions.

**Benefits :** Simple and efficient, improved performance, scalability, manageability

**2. Hash partitioning:**

Hash partitioning is a type of horizontal partitioning in which the *data is divided into multiple partitions based on a hash function.* A hash function is a mathematical formula that takes a value as input and produces a fixed-length output, called a *hash or a bucket.*

To hash partition a table, the database first applies the hash function to a column or a combination of columns in the table, called the *partition key.* The resulting hash value is then used to determine which partition should hold a particular piece of data.

**3. Range partitioning:**

This type of partitioning ***divides the data into ranges, based on a specific column value. Each range being stored in a separate table.***

For example, you could range partition a customer table by customer ID, with each range containing a specific subset of customer IDs.

Example

Suppose we have a table of customer data, with the following columns:

***customer\_id***

***first\_name***

***last\_name***

***email***

***address***

***phone\_number***

We could range partition this table by customer ID, with each partition containing a range of customer IDs.

For example, we could create three partitions, one for customer IDs 0-9999, one for customer IDs 10000-19999, and one for customer IDs 20000 and above.

When a new customer record is inserted into the table, the database would determine which partition the record belongs to by comparing the customer ID value to the range of values for each partition. For example, a customer record with the customer ID 12345 would be assigned to the partition for customer IDs 0-9999.

## 2. VERTICAL PARTITIONING

Vertical partitioning divides the data into multiple tables, with each table containing a subset of the columns. This type of partitioning is ***less common than horizontal partitioning***, but it can be useful for certain types of queries, such as queries that only need to access a subset of the columns in a table.

**There are two main types of vertical partitioning:**

**1. Column partitioning:** This involves dividing the data into multiple tables, with ***each table containing a different set of columns.***

**2. Table partitioning:** This involves dividing the data into multiple tables, with each table ***containing a different version of the same table.***

## 4. TYPES OF SKEWS AND HANDLING [UNIT 3]

**Skew** in the context of databases refers to an ***imbalance or uneven distribution of data or workload.*** When skew occurs, ***some data elements or processing units are significantly more***

**loaded** than others. Skew can negatively impact the performance and efficiency of a database system.

**There are two main types of skews:**

1. **Attribute Skew (Data Skew):** Attribute skew occurs when the attribute of a database table is not evenly distributed. **Some values occur much more frequently than others**, leading to uneven data distribution.
  - **Example:** In a customer database, the "State" column might exhibit attribute skew if one state (e.g., California) has significantly more customer records than others.
2. **Partition Skew (Workload Skew):** Partition skew happens **when the workload or data access patterns are not evenly distributed** across partitions or processing units. This means that some partitions are handling a much higher load of queries or transactions than others.
  - **Example:** In a distributed database with multiple nodes, partition skew occurs when one node receives many more queries or transactions than the others, leading to overutilization of resources.

### 1. Handling Skew in Range Partitioning:

**To create a balanced partitioning vector (assuming partitioning attribute forms a key of the relation):**

1. **Sort the relation** on the **partitioning attribute**.
2. **Construct the partition vector** by **scanning the relation in sorted order** as follows.
  - After every  $1/n$  th of the relation has been read, the value of the partitioning attribute of the next tuple is added to the partition vector.
3.  **$n$  denotes the number of partitions to be constructed.**
4. Duplicate entries or imbalances can result if duplicates are present in partitioning attributes.

### 2. Handling Skew Using Histogram:

Histograms are **statistical representations of data distribution**. To handle skew using histograms:

The **database optimizer can use histogram statistics to make better query execution plans**. For example, it can choose different **join strategies or access paths** based on the distribution of data, potentially mitigating the effects of skew.

### 3. Handling Skew Using Virtual Processors:

Virtual processors, also known as query virtualization, involve dynamically allocating resources to queries based on their complexity and requirements. To handle skew using virtual processors:

- **Resource Allocation:** Allocate more processing resources (CPU, memory) to queries or transactions that involve skewed data. This helps ensure that skewed queries complete faster, preventing them from blocking other queries.
- **Dynamic Load Balancing:** Implement dynamic load balancing mechanisms that distribute query workloads across available virtual processors based on the query complexity and the presence of skew.
- **Query Prioritization:** Prioritize and schedule queries with skew so that they receive more resources and are processed efficiently without adversely affecting other queries.

Handling skew in a database is essential for maintaining good performance and preventing bottlenecks. The specific approach used will depend on the type of skew, the database system's capabilities, and the characteristics of the data and workload.

## 5. EXPLAIN THE TERM INTERQUERY & INTRAQUERY PARALLELISM, PARALLEL SORT & JOIN. [UNIT 3]

In the context of parallel databases, several techniques and concepts are used to improve query performance by taking advantage of parallel processing capabilities. Two fundamental concepts related to parallelism in databases are *"interquery parallelism"* and *"intraquery parallelism"*.

1. **Interquery Parallelism:** Interquery parallelism refers to the *parallel execution of multiple queries or transactions simultaneously*. In a parallel database system, multiple queries from different users or applications can be processed concurrently, with each query running on its own set of resources, such as *CPU cores, memory, and storage*. This enables the database system to serve multiple users concurrently, *improving overall throughput* and *reducing response times* for individual queries.

Interquery parallelism is *essential for multi-user database systems*, as it allows concurrent access to the database without significant contention for resources.

2. **Intraquery Parallelism:** Intraquery parallelism, also known as *query parallelism*, refers to the *parallel execution of a single query or transaction*. When a complex query is submitted to the database, it can be *broken down into multiple smaller subqueries* that can be *executed concurrently on different processing units*. Each subquery works on a subset of the data, and their results are combined to produce the final result of the overall query.

Intraquery parallelism is particularly useful for *speeding up the execution of complex analytical queries*, which often involve *large datasets and complex computations*. Parallel execution can significantly reduce the query's response time by dividing the workload among multiple resources.

." Additionally, "parallel sort" and "parallel join" are techniques used to optimize specific database operations.

1. **Parallel Sort:** Parallel sorting is a technique *used to speed up the sorting of large datasets* in a parallel database environment. Sorting is a common operation in many database queries, such as *ORDER BY clauses* or *building indexes*. In a parallel sort, the *database system divides the data into smaller partitions and sorts each partition concurrently using multiple processing units*. Once the partitions are sorted, they can be merged to produce the final sorted result.

2. **Parallel Join:** Parallel join is a technique used to optimize join operations in database queries, where data from two or more tables are combined based on a specified condition.



***Join operations can be resource-intensive, especially when dealing with large tables.***

In a parallel join, the database system divides the join operation into smaller tasks that can be executed concurrently on multiple processing units. Each task processes a subset of the data and produces intermediate results, which are then combined to produce the final join result.

In summary, interquery parallelism and intraquery parallelism enable concurrent processing of multiple queries and parallel execution of a single query, respectively, in a parallel database system. Parallel sort and parallel join are specific techniques used to optimize sorting and join operations, improving the overall performance of database queries, especially when dealing with large datasets.

## **6. EXPLAIN PARALLEL SORT IN DETAIL WITH A SUITABLE EXAMPLE/ GIVE DETAILS OF PARALLEL SORT. [ UNIT3 ]**

Parallel sort is a technique used in database systems to efficiently sort large sets of data by dividing the sorting task among multiple processors. This approach takes advantage of intraoperation parallelism, where operations on relations are parallelized across different subsets of the data.

### **Two Approaches for Parallel Sort:**

#### **1. Range-Partitioning Sort:**

- Step 1: Tuples are redistributed to processors based on a range-partition strategy.
- Step 2: Each processor independently sorts its partition.
- Merge: Since tuples are range-partitioned, the final merge is straightforward, and each processor's sorted data can be concatenated to obtain the fully sorted relation.
- Example: Suppose you have a dataset of employee records to be sorted by their salary. The range-partitioning sort would distribute records to processors based on salary ranges, and each processor would independently sort its subset.

#### **2. Parallel External Sort-Merge:**

- Step 1: Each processor sorts its local data independently.
- Step 2: The system merges sorted runs on each processor to obtain the final sorted output.
- Merge Process:
  - The system range-partitions the sorted partitions across processors, sending tuples in sorted order.
  - Each processor merges the received sorted streams to create a single sorted run.
  - The system concatenates the sorted runs from all processors to produce the final result.

#### **Considerations:**

Efficient range partitioning or parallel external sort-merge requires a good partitioning strategy.

Skew issues, where data distribution is uneven, can be addressed using techniques like virtual processor partitioning.

### Example:

Sure, let's consider a small example with a simplified `employee` table:

```
sql Copy code

CREATE TABLE employee (
  employee_id INT,
  employee_name VARCHAR(50),
  salary INT
);

INSERT INTO employee VALUES
  (1, 'Alice', 60000),
  (2, 'Bob', 45000),
  (3, 'Charlie', 75000),
  (4, 'David', 50000),
  (5, 'Eva', 70000);
```

Now, let's perform a parallel sort on the `employee` table based on the `salary` column.

#### 1. Range-Partitioning Sort:

```
sql Copy code

-- Step 1: Range partitioning
-- Each processor sorts its partition independently

-- Processor 1
SELECT * INTO TEMPORARY TABLE temp_sort_1
FROM employee
WHERE salary BETWEEN 0 AND 50000;

-- Processor 2
SELECT * INTO TEMPORARY TABLE temp_sort_2
FROM employee
WHERE salary BETWEEN 50001 AND 100000;

-- Step 2: Concatenate sorted partitions
SELECT * FROM temp_sort_1
UNION ALL
SELECT * FROM temp_sort_2
ORDER BY salary;
```

#### 2. Parallel External Sort-Merge:

```
sql Copy code

-- Step 1: Each processor sorts its local data

-- Processor 1
SELECT * INTO TEMPORARY TABLE temp_sort_1
FROM employee
WHERE salary BETWEEN 0 AND 50000
ORDER BY salary;

-- Processor 2
SELECT * INTO TEMPORARY TABLE temp_sort_2
FROM employee
WHERE salary BETWEEN 50001 AND 100000
ORDER BY salary;

-- Step 2: Parallel merge of sorted runs
-- Merge sorted runs from all processors

-- Final merge
SELECT * FROM temp_sort_1
UNION ALL
SELECT * FROM temp_sort_2
ORDER BY salary;
```

In this example, we have a small `employee` table with different salary values. The parallel sort is simulated by partitioning the data based on salary ranges and then either concatenating sorted partitions or merging sorted runs in parallel.

## 7. EXPLAIN THE CONCEPT OF PARALLEL JOIN IN DETAIL. [UNIT 3]

Parallel join is a concept in parallel database systems where the traditional join operation, which involves testing pairs of tuples from two relations to check if they satisfy a specified condition, is parallelized to improve performance.

There are different techniques for achieving parallel join, and they can be broadly categorized into partitioned join, fragment-and-replicate join, and parallelized versions of specific join algorithms.

### 1. Partitioned Join:

- *Idea:* Partition both input relations across multiple processors and compute the join locally at each processor.
- *Methods:* Range partitioning or hash partitioning based on join attributes.
- *Requirements:* Suitable for Equi-joins where the join condition involves equality. Relations must be partitioned using the same partitioning function.
- *Example:* If there are  $n$  processors and two relations  $r$  and  $s$ , they are partitioned into  $r_0, r_1, \dots, r_{n-1}$  and  $s_0, s_1, \dots, s_{n-1}$ . Local join is computed at each processor  $P_i$  for  $r_i$  and  $s_i$ .

### 2. Fragment-and-Replicate Join:

- *Idea:* Suitable for joins with inequality conditions. It involves partitioning one relation and replicating the other across processors.
- *Methods:* Asymmetric fragment and replicate (replicate one relation and partition the other) or general fragment and replicate (partition both relations independently).
- *Usage:* When partitioning is not applicable, especially when dealing with inequality conditions.
- *Example:* Replicate one relation (e.g.,  $s$ ) across all processors and partition the other relation (e.g.,  $r$ ) into multiple partitions. Each processor computes the join of its local partition of  $r$  with the replicated  $s$ .

### 3. Parallelized Join Algorithms:

- *Hash Join:* Parallelizing the partitioned hash join algorithm, which involves distributing tuples based on hash functions and locally computing the join.
- *Nested-Loop Join:* Replicating a smaller relation and using asymmetric fragment-and-replicate parallelization for nested-loop joins.

### Optimizations:

- *Buffering:* Optimizations such as buffering tuples in memory to reduce I/O costs.
- *Synchronization:* Ensuring synchronization between the replication of one relation and the join operation to manage in-memory buffers efficiently.

Parallel join techniques aim to distribute the workload among processors, reducing the time required to perform join operations on large datasets and improving overall system performance in a parallel database environment.

Suppose we have two relations: Employees and Departments. We want to perform a parallel join on the "DepartmentID" attribute, which is common to both relations.

#### Employees Table:

```
sql Copy code

CREATE TABLE Employees (
    EmployeeID INT,
    EmployeeName VARCHAR(255),
    DepartmentID INT
);

INSERT INTO Employees VALUES
(1, 'John Doe', 101),
(2, 'Jane Smith', 102),
(3, 'Bob Johnson', 101),
(4, 'Alice Williams', 103);
```

#### Departments Table:

```
sql Copy code

CREATE TABLE Departments (
    DepartmentID INT,
    DepartmentName VARCHAR(255)
);

INSERT INTO Departments VALUES
(101, 'HR'),
(102, 'Marketing'),
(103, 'Finance');
```

Now, let's perform a parallel join using the SQL `JOIN` operation. In a real parallel database system, the parallelism would be handled automatically by the system based on the data distribution across processors. In this simplified example, we'll use a basic `JOIN` statement.

#### Parallel Join Example:

```
sql Copy code

SELECT Employees.EmployeeID, Employees.EmployeeName, Employees.DepartmentID,
FROM Employees
JOIN Departments ON Employees.DepartmentID = Departments.DepartmentID;
```

This SQL query retrieves information about employees and their corresponding departments by joining the Employees and Departments tables on the DepartmentID attribute. In a parallel database system, the actual parallelization would be handled behind the scenes, and the query optimizer would determine the most efficient way to distribute and process the data across multiple processors.

## 8. EXPLAIN TERMS. [UNIT 3]

### 1. Pipeline parallelism

Pipeline parallelism is a technique for improving the performance of parallel applications by dividing the computation into a sequence of stages and executing them simultaneously on different processors. This is similar to how an assembly line works, where each stage of the line performs a different task on the product as it moves through.

In pipeline parallelism, the stages are typically arranged in a linear sequence, with the output of one stage being the input to the next stage. The stages are also typically designed to be independent of each other, so that they can be executed simultaneously without interfering with each other.

Pipeline parallelism can be used to improve the performance of a wide variety of parallel applications, including machine learning algorithms, scientific computing applications, and data processing applications.

#### Benefits of pipeline parallelism

Pipeline parallelism provides a number of benefits, including:

- **Improved performance:** Pipeline parallelism can significantly improve the performance of parallel applications by allowing different parts of the computation to be executed simultaneously.
- **Increased scalability:** Pipeline parallelism can help to increase the scalability of parallel applications by allowing the computation to be distributed across multiple processors.
- **Reduced memory usage:** Pipeline parallelism can reduce the memory usage of parallel applications by only storing the data that is needed for the current stage of the computation.

#### Example

Suppose we have a parallel machine learning algorithm that trains a neural network. The algorithm can be divided into the following stages:

1. Load the training data.
2. Preprocess the training data.
3. Train the neural network.
4. Evaluate the neural network.

We can use pipeline parallelism to execute these stages simultaneously on different processors. For example, we could load the training data on one processor, preprocess the training data on another processor, train the neural network on a third processor, and evaluate the neural network on a fourth processor.

By executing the stages of the algorithm simultaneously, we can significantly improve the performance of the algorithm.

### 2. Parallel aggregate

Parallel aggregate in a distributed database is the process of aggregating data from multiple nodes in the cluster in parallel. This can be done by performing partial aggregations on each node and then combining the partial aggregations on a single node. Parallel aggregate can significantly improve the performance of queries that involve aggregation of large datasets.

**Here is a simpler explanation of parallel aggregate with an example:**

Imagine you have a very large table of sales data, and you want to calculate the total sales for each product category. This is a common aggregation query, but it can be very slow if you have a large amount of data.

One way to speed up this query is to use parallel aggregate. Parallel aggregate works by dividing the data into smaller chunks and processing each chunk independently on a separate node in the cluster. This allows the cluster to aggregate the data much more quickly than if it were to aggregate the data sequentially.

Here is an example of how parallel aggregate might be used to calculate the total sales for each product category:

1. The cluster divides the sales table into smaller chunks, based on product category.
2. Each node in the cluster aggregates the sales data for its assigned chunk.
3. The cluster then combines the partial aggregates from each node to calculate the total sales for each product category.

This approach can significantly improve the performance of aggregation queries on large datasets.

**Here is a simple example of a parallel aggregate query in SQL:**

```
SELECT product_category, SUM(sales)
FROM sales_table
GROUP BY product_category
PARALLEL BY product_category;
```

This query will calculate the total sales for each product category, using parallel aggregate.

Parallel aggregate is a powerful technique that can be used to improve the performance of aggregation queries on large datasets in a distributed database.

**ARRAY EXAMPLE IN SQL:**

```
CREATE TYPE phone_number_array AS ARRAY[VARCHAR(255)];

CREATE TABLE employees (
  employee_id INT NOT NULL AUTO_INCREMENT,
  name VARCHAR(255) NOT NULL,
  phone_numbers phone_number_array,
  PRIMARY KEY (employee_id)
);

INSERT INTO employees (name, phone_numbers) VALUES
  ('John Smith', ARRAY['(555) 555-5555', '(415) 555-5555']),
  ('Jane Doe', ARRAY['(123) 456-7890', '(987) 654-3210']);
```

## MULTISET EXAMPLE IN SQL:

```
CREATE TYPE phone_number_multiset AS MULTISET[VARCHAR(255)];

CREATE TABLE employees (
  employee_id INT NOT NULL AUTO_INCREMENT,
  name VARCHAR(255) NOT NULL,
  phone_numbers phone_number_multiset,
  PRIMARY KEY (employee_id)
);

INSERT INTO employees (name, phone_numbers) VALUES
  ('John Smith', MULTISET['(555) 555-5555', '(415) 555-5555']),
  ('Jane Doe', MULTISET['(123) 456-7890', '(987) 654-3210']);
```

## WITH A SUITABLE EXAMPLE EXPLAIN SQL FUNCTIONS AND PROCEDURES.

### SQL Functions:

SQL functions are special operations or computations that can be performed on the data stored in a database. They take one or more parameters as input, perform a specific action, and return the result. Functions can be used in SQL queries to manipulate data or perform calculations. There are various types of SQL functions, such as string functions, numeric functions, date functions, and aggregate functions.

### Example:

Let's consider the **UPPER** function. The **UPPER** function is used to convert a string to uppercase. The theoretical concept is that it transforms all characters in a given string to their uppercase equivalents.

```
-- Example Usage

SELECT UPPER('hello') AS uppercase_result;
```

In this example, the **UPPER** function takes the string 'hello' as input and returns 'HELLO' as the result. The function performs the conversion of lowercase characters to uppercase, illustrating the theoretical concept.

## SQL Procedures:

SQL procedures are sets of SQL statements grouped together to perform a specific task or a series of tasks. They are similar to functions but may not necessarily return a value. Procedures are useful for encapsulating a sequence of SQL statements that need to be executed as a single unit. They can include control flow statements, loops, and error handling.

### Example:

Let's create a simple SQL procedure named **IncreaseSalary** that increases the salary of an employee by a specified percentage.

```
-- Example Usage
DELIMITER //
CREATE PROCEDURE IncreaseSalary(IN employee_id INT, IN percentage INT)
BEGIN
    UPDATE employees
    SET salary = salary + (salary * percentage / 100)
    WHERE employee_id = employee_id;
END //
DELIMITER ;

-- Execute the procedure
CALL IncreaseSalary(101, 10);
```

In this example, the **IncreaseSalary** procedure theoretically takes an **employee\_id** and a **percentage** as input, then updates the salary of the specified employee by increasing it by the given percentage. The **DELIMITER** statements are used to define the procedure, and the procedure is executed using the **CALL** statement.

## 1. EXPLAIN TUNING OF SCHEMA AND TUNING OF INDICES IN DBMS PERFORMANCE TUNING. [UNIT 4]

### Tuning of Schema:

#### 1. Vertical Partitioning:

- Splitting relations based on common access patterns within normal form constraints.
- Example: For a "course" relation, attributes like credits can be separated for better performance.

#### 2. Column Store Approach:



- Storing each attribute in a separate file.
- Effective for data warehouse applications.

### 3. **Denormalized Relations:**

- Join information is denormalized to reduce query time.
- Requires careful consistency maintenance.

### 4. **Materialized Views:**

- Offering similar benefits to denormalized relations.
- Helps in maintaining consistency.

### 5. **Clustered Records:**

- Clustering records that match in a join on the same disk page.
- Utilizes clustered file organizations.

## **Tuning of Indices:**

### 1. **Appropriate Indices:**

- Creating indices on relations to speed up queries.
- Removing unnecessary indices if updates are a bottleneck.

### 2. **Index Types:**

- Choosing index types based on the workload.
- B-tree indices for range queries, hash indices for other cases.

### 3. **Clustered Index:**

- Only one index on a relation can be made clustered.
- Typically, the index benefiting the most queries and updates should be clustered.

### 4. **Tuning Wizards:**

- DBMS provides tuning wizards that analyse past queries and update history (workload).
- These tools suggest indices based on workload estimates for improved execution time.

In summary, tuning the schema involves organizing relations for optimal access patterns, and tuning indices focuses on creating, removing, and selecting index types based on workload characteristics. These optimizations aim to enhance the overall performance of a DBMS.

## **2. EXPLAIN TPC BENCHMARK OF PERFORMANCE. [UNIT 4]**

The Transaction Processing Performance Council (TPC) defines a series of benchmarks to measure the performance of database systems. Here's an explanation of the TPC benchmarks:

### 1. Benchmark Overview:

- Purpose: Evaluate the performance of database systems.
- Metrics: Focus on throughput, measured in transactions per second (TPS).
- Response Time: Systems must meet certain response time bounds.

### 2. Performance Metrics:

- Throughput: Measured in transactions per second (TPS).
- Price per TPS: Evaluates system cost-effectiveness.
- External Audit: Systems must undergo external audits to ensure adherence to benchmark definitions, including support for ACID properties.

### 3. Benchmark Types:

- TPC-A: Simulates a bank application for cash withdrawal and deposit.
- TPC-B: Tests core database system performance, excluding user-related aspects.
- TPC-C: Models complex order-entry environments, widely used for online transaction processing (OLTP) systems.
- TPC-E: Aimed at OLTP systems, models a brokerage firm's interactions with customers and financial markets.
- TPC-D: Assesses performance on decision-support queries, using a sales/distribution application.
- TPC-H (Ad Hoc): A refinement of TPC-D with 22 queries, focusing on ad hoc querying. Prohibits materialized views and redundant information.
- TPC-W (Web Commerce): Models end-to-end performance of Web sites with static and dynamic content. Focuses on Web interactions per second (WIPS) and price per WIPS.

### 4. TPC-D and Materialized Views:

- TPC-D queries can benefit from materialized views, especially in applications with known, repeated queries.
- Users realized queries could be sped up, but overhead of maintaining materialized views must be considered.

### 5. TPC-H Benchmark:

- A refinement of TPC-D with 22 queries, inserts, deletes, and updates.
- Prohibits materialized views and redundant information.

- Measures performance through power and throughput tests, with a composite metric reflecting queries per hour.
- Defines a composite price/performance metric.

#### **6. TPC-W Web Commerce Benchmark:**

- Models Web sites with static and dynamic content.
- Allows caching of dynamic content for improved performance.
- Measures Web interactions per second (WIPS) and price per WIPS.

In summary, TPC benchmarks provide standardized ways to assess and compare the performance of database systems across various scenarios, ensuring fair and meaningful evaluations.

The benchmarks cover a range of application types, from transaction processing to decision support and web commerce.

### **3. DISCUSS THE CONCEPT OF E-COMMERCE. [UNIT 4]**

E-commerce, short for electronic commerce, refers to the buying and selling of goods and services over the internet. This digital form of commerce has transformed the way businesses operate and how consumers shop. The concept of e-commerce encompasses a wide range of online activities, and it has become a vital aspect of the global economy. Here are key aspects to consider when discussing the concept of e-commerce:

#### **1. Online Transactions:**

- **Buying and Selling:** E-commerce involves the online exchange of goods and services. Businesses can sell their products or services to consumers, and consumers can make purchases from the comfort of their homes or through mobile devices.
- 

#### **2. Types of E-commerce:**

- **B2C (Business-to-Consumer):** In B2C e-commerce, businesses sell products and services directly to consumers. Examples include online retailers, such as Amazon and eBay.
- **B2B (Business-to-Business):** B2B e-commerce involves transactions between businesses. Companies purchase goods or services from other companies through online platforms.
- **C2C (Consumer-to-Consumer):** C2C e-commerce facilitates transactions between consumers. Online marketplaces enable individuals to buy and sell directly to each other. Examples include platforms like Etsy and Craigslist.

#### **3. Advantages of E-commerce:**

- **Global Reach:** E-commerce provides businesses with a global reach, allowing them to reach customers beyond geographical boundaries.
- **Convenience:** Consumers can shop at any time, from anywhere, providing a high level of convenience.
- **Cost Savings:** Online businesses often have lower operational costs than traditional brick-and-mortar stores, leading to potential cost savings.

#### **4. E-commerce Components:**

- **Online Stores:** Businesses set up digital storefronts where customers can browse and purchase products.
- **Payment Systems:** Secure online payment methods, such as credit cards, digital wallets, and other payment gateways, enable transactions.
- **Shopping Carts:** Virtual shopping carts allow users to add items, review selections, and proceed to checkout.
- **Security Measures:** E-commerce platforms implement robust security measures to protect customer data and ensure secure transactions.

#### **5. Challenges and Considerations:**

- **Security Concerns:** E-commerce faces challenges related to cybersecurity, including the protection of sensitive customer information.
- **Logistics and Fulfilment:** Efficient order fulfilment and delivery are crucial components, and businesses need reliable logistics for timely shipments.
- **Competition:** The digital marketplace is highly competitive, requiring businesses to implement effective marketing and customer engagement strategies.

#### **6. Mobile Commerce (M-commerce):**

- With the widespread use of smartphones and tablets, M-commerce refers to e-commerce conducted through mobile devices. This includes mobile apps, mobile-optimized websites, and mobile payment systems.

#### **7. Future Trends:**

- **AI and Personalization:** Artificial intelligence is being used to enhance user experiences through personalized recommendations.
- **Voice Commerce:** The rise of voice-activated devices has led to the emergence of voice commerce, where users can make purchases using voice commands.

E-commerce continues to evolve with technological advancements, shaping the way businesses connect with consumers and how consumers access products and services. The ongoing integration of digital technologies further contributes to the growth and transformation of the e-commerce landscape.

## 1. WHAT IS DECISION SUPPORT SYSTEM? EXPLAIN.[UNIT 5]

A Decision Support System (DSS) is a category of database applications designed to extract high-level information from detailed transaction data stored in transaction-processing systems. The primary goal of a Decision Support System is to aid managers in making informed decisions based on the analysis of accumulated data.

Here's the list of key points related to DSS:

1. **Data Analysis:** DSS allows users to analyse large amounts of data from various sources to gain insights and make informed decisions. It provides tools for examining, interpreting, and visualizing data to identify patterns and trends.
2. **OLAP (Online Analytical Processing):** DSS often incorporates OLAP, which enables users to analyse multidimensional data from different perspectives. This allows for in-depth exploration of data to support decision-making processes.

3. **Extended Aggregation Features in SQL:** DSS leverages SQL with extended aggregation features to perform complex calculations and analysis on large datasets. This includes advanced functions for windowing, ranking, and aggregating data to derive meaningful conclusions.
  - **Windowing and Ranking:** These features in SQL allow users to perform calculations over a set of rows related to the current row, enabling comparative analysis and ranking of data based on specific criteria.
4. **Data Mining:** DSS integrates data mining capabilities, allowing users to discover patterns, correlations, and anomalies within the data. Data mining techniques help uncover valuable insights that can guide decision-making processes.
5. **Data Warehousing:** DSS often relies on data warehousing, which involves the collection, storage, and organization of data from various sources into a centralized repository. This allows for easy access to integrated data for analysis.

In simple terms, a Decision Support System is a computer-based tool that helps people make better decisions by providing tools for analysing data, exploring trends and patterns, performing advanced calculations, discovering insights through data mining, and accessing a centralized repository of organized data for decision-making purposes.

## 2. EXPLAIN DATA WAREHOUSE ARCHITECTURE.[UNIT 5]

Data warehousing is like a super-organized library for big companies with lots of data. Imagine these companies having stores all over, each generating heaps of information. Data might be scattered, like one place tracking customer complaints and another dealing with manufacturing problems. Also, they might buy data from outside, like credit scores for customers. When bosses need to make decisions, they want info from all these places, not just one.

That's where data warehouses come in. Think of them as a central hub that collects data from everywhere and organizes it neatly. It's like putting books from different branches of a library into one big library. This way, decision-makers can easily find what they need without dealing with messy individual sources.




Figure 20.1 Data-warehouse architecture.

The process involves some steps:

1. **Gathering Data:** There are different ways to collect data, either by regularly asking for updates or by getting them as they happen.
2. **Schema and Integration:** Since different places might use different ways to organize data, the warehouse creates a common system (schema) to blend everything together.
3. **Cleaning and Transforming:** Sometimes data might be messy, with misspelled names or wrong addresses. The warehouse fixes these issues and transforms data to fit the common system.
4. **Updates and Summaries:** Changes in the original data need to be reflected in the warehouse. Also, instead of storing every little detail, the warehouse might keep summaries for quick access.

This whole process is called ETL: Extract, Transform, and Load. It's like gathering books (extract), organizing them (transform), and putting them on shelves (load).

The key is that a data warehouse gives decision-makers a clean, historical, and one-stop-shop for all the data they need. It's like a super-efficient library that makes finding information a breeze.

### 3. WHAT IS DATA MINING? DESCRIBE THE APPLICATIONS OF DATA MINING.[UNIT 5]

#### Data Mining:

Data mining is the process of discovering patterns, trends, correlations, or useful information from large sets of data. It involves using various techniques from machine learning, statistics, and database systems to analyse and interpret data. The goal is to uncover hidden patterns and knowledge, which can be valuable for making informed decisions in business, research, and other domains.

#### Applications of Data Mining:

##### 1. Prediction based on past history:

- Predicting credit card risk based on attributes and past history.
- Predicting customer brand loyalty.
- Predicting customer response to "junk mail."
- Predicting fraudulent phone calling card usage.

##### Mechanisms:

- **Classification:** Predicting the class (e.g., good or bad credit risk, loyal or likely to switch loyalty).
- **Regression formulae:** Predicting a numeric result based on parameters.

##### 2. Descriptive Patterns:

- **Associations:** Discovering patterns of items bought together, like suggesting additional books based on a customer's purchase.
- **Clusters:** Identifying clusters of related cases, like detecting a cluster of typhoid cases around a contaminated well.

### 3. Applications in Detecting Causation:

- Identifying associations as a preliminary step in detecting causation (e.g., chemical exposure and cancer, or medicine and cardiac problems).

In summary, these applications demonstrate the diverse utility of data mining techniques in making predictions, uncovering patterns, and identifying relationships within large datasets. These insights can be invaluable for decision-making, risk assessment, and improving various aspects of business and healthcare.

## 4. EXPLAIN CLASSIFICATION AND DECISION TREE CLASSIFIER.[UNIT 5]

### Classification Overview:

Classification, a pivotal aspect of data mining, revolves around predicting the class to which a new item belongs based on past instances and their associated classes, known as the training instances. In this context, a class can represent categories like creditworthiness levels or any distinct groups relevant to the application.

### Decision-Tree Classifiers:

Among various classification techniques, decision-tree classifiers stand out. These classifiers employ a tree structure where each leaf node corresponds to a class, and internal nodes host predicates or functions. To classify a new instance, one traverses the tree from the root to a leaf, evaluating predicates along the way. An example might involve deciding the credit risk of a person based on attributes like education level and income.

### Building Decision-Tree Classifiers:

The process of constructing a decision-tree classifier initiates with a training set—a sample of data with known class labels. **A greedy algorithm, working recursively, starts from the root, associating all training instances with it. Nodes become leaves if the instances share the same class.** Otherwise, a **partitioning attribute** is chosen to create child nodes, and the process repeats.

1. **Greedy Algorithm:** The process begins with a training set, associating all instances with the root.
2. **Node Expansion:** Nodes become leaves if instances share the same class or are too few; otherwise, a partitioning attribute is chosen.
3. **Recursive Process:** The algorithm continues recursively, building the tree until a stopping criterion is met.

### Best Splits and Purity Measures:



The efficacy of a split is gauged by measures like Gini impurity or entropy. Gini measures how often a randomly selected element would be incorrectly labelled, while entropy comes from information theory. The best split for an attribute is the one maximizing the information gain ratio, indicating improved classification purity.

#### **Gini Purity:**

$$\text{Gini}(S) = 1 - \sum_{i=1}^k p_i^2$$

#### **Entropy:**

$$\text{Entropy}(S) = - \sum_{i=1}^k p_i \log_2(p_i)$$

#### **Other Types of Classifiers:**

1. Logistic Regression:
2. Decision Trees:
3. Random Forest:
4. Support Vector Machines (SVM):
5. K-Nearest Neighbours (KNN):

In essence, classification and decision-tree classifiers play a crucial role in extracting valuable insights from data, aiding predictions, and automating decision-making processes across various domains. These techniques empower systems to categorize, predict, and derive meaningful patterns from datasets, facilitating informed decision-making.

### **5. DISCUSS PAGE RANKING AND POPULARITY RANKING IN INFORMATION RETRIEVAL. [UNIT 5]**

**Ranking** is the process of arranging items in a specific order based on their relevance, importance, or some predefined criteria. In various domains, including information retrieval, ranking is crucial for presenting users with the most relevant and valuable items first.

#### **Popularity Ranking:**

##### **Definition:**

- Popularity ranking, also known as prestige ranking, aims to find and rank pages that are popular higher than others containing the specified keywords.
- The popularity of a page is influenced by the number of hyperlinks pointing to it.

##### **Key Points:**

1. **Importance of Hyperlinks:** Early search engines relied solely on TF-IDF-based relevance measures, but this had limitations on the web due to the abundance of pages with all query keywords.
2. **Challenges:** Determining the popularity of a page is challenging. Access frequency data is often private or unreliable if provided by the site.

3. **Using Hyperlinks:** Hyperlinks are used as a measure of popularity. Bookmark files and links from other sites are analysed to infer a page's popularity.
4. **Link to Page or Site:** The popularity can be associated with either individual pages or entire sites. The association with sites can mitigate issues with external links often pointing only to the root page.

#### Critique:

- The passage notes challenges in accurately measuring popularity, especially when relying on the number of external links.

#### PageRank:

##### Definition:

- PageRank is a measure of a page's popularity based on the popularity of pages linking to it.
- Introduced by Google, it significantly improved search engine ranking by considering link structure.
- Then the PageRank  $P[j]$  for each page  $j$  can be defined as:

$$P[j] = \delta/N + (1 - \delta) * \sum_{i=1}^N (T[i, j] * P[i])$$

#### Key Points:

1. **Random Walk Model:** PageRank can be understood using a random walk model where a person traverses web pages randomly, either jumping to a random page or following links.
2. **Link Influence:** Pages with more incoming links are more likely to be visited, contributing to a higher PageRank.
3. **Linear Equations:** PageRank can be defined by a set of linear equations, where the probability of a random walk being a jump is a key parameter.
4. **Iterative Technique:** The equations are solved using an iterative technique, starting with initial PageRank values and updating them until convergence.

#### Critique:

- PageRank overcomes some of the challenges of traditional relevance measures by incorporating the link structure, providing more accurate results.

#### Summary:

- **Popularity Ranking:** Primarily relies on the number of hyperlinks to a page, considering both external links and links from other pages within the site.
- **PageRank:** Incorporates a more sophisticated approach by considering the structure of the entire link network, providing improved relevance in web search.

These approaches highlight the evolution from traditional relevance measures to more advanced methods that leverage hyperlink information for ranking web pages.

## **6. EXPLAIN CRAWLING AND INDEXING IN WEB DATABASE. [UNIT 5]**

### **Crawling**

Crawling is the process of discovering new web pages. A web crawler, also known as a spider or robot, is a software program that visits web pages and follows the links on those pages to discover new pages. Crawlers start with a list of known web pages and then recursively follow the links on those pages until they have discovered all of the pages that they can reach.

### **Indexing**

Indexing is the process of storing and organizing the content of web pages so that it can be searched efficiently. When a crawler visits a web page, it indexes the page's content by extracting the page's title, meta tags, and body text. The crawler also indexes the links on the page so that it can follow those links and discover new pages.

### **Crawling and indexing in a web database:**

In a web database, crawling and indexing are used to create a searchable database of web pages. The database is stored on a server and can be accessed by users to search for web pages.

### **How it works:**

The following is a simplified explanation of how crawling and indexing work in a web database:

1. The web crawler starts with a list of known web pages.
2. The crawler visits each web page in the list and extracts the page's content.
3. The crawler indexes the page's content and stores it in the database.
4. The crawler follows the links on the page and discovers new pages.
5. The crawler adds the new pages to the list of known pages and repeats steps 2-4.

This process continues until the crawler has discovered and indexed all of the web pages that it can reach.

### **Benefits of crawling and indexing:**

Crawling and indexing have a number of benefits, including:


- Search engine optimization (SEO): Crawling and indexing are essential for SEO, which is the process of optimizing a website so that it ranks higher in search engine results pages (SERPs).
- User experience: Crawling and indexing make it possible for users to search for web pages efficiently and find the information they are looking for quickly.

- Data analysis: Crawled and indexed web pages can be analysed to gain insights into user behaviour, trends, and other topics.

## 1. WITH A NEAT DIAGRAM EXPLAIN TP-MONITOR ARCHITECTURES.

### Transaction Processing Monitor (TP Monitor)

A TP monitor is a software application that coordinates and manages transactional activities in a distributed computing environment. It ensures the efficient and reliable execution of transactions across multiple interconnected systems.



## Types of TP Monitor Architectures:

### 1. Process per Client Model:

- In this model, each client is associated with a dedicated server process, which handles the client's requests and communication.
- It involves high memory requirements because per client dedicated server process is initiated.
- The multitasking nature of this model leads to high CPU overhead for context switching between processes.

### 2. Single Process Model:

- In this architecture, all remote clients connect to a single server process, commonly used in client-server environments.
- The server process is multi-threaded, resulting in lower cost for context switching.
- However, there is no protection between applications, and it is not well-suited for parallel or distributed databases.

### 3. Many-Server Single-Router Model:

- Multiple application server processes access a common database, while clients communicate with the application through a single communication process that routes requests.
- It involves independent and multithreaded server processes for multiple applications.
- This model is suitable for running on parallel or distributed databases.

### 4. Many-Server Many-Router Model:

- In this model, multiple processes communicate with clients, and client communication processes interact with router processes to route requests to the appropriate server.

- A controller process starts up and supervises other processes in this architecture.

These TP monitor architectures are designed to manage and coordinate transactional activities in a distributed environment, ensuring efficient communication between clients and server processes.

## 2. EXPLAIN WITH EXAMPLE AND NEAT DIAGRAM, THE TRANSACTION WORKFLOWS.

Transactional workflows are intricate processes that involve the coordinated execution of multiple tasks by different entities, commonly utilized in scenarios where interactions between various systems or involving human interventions, such as loan processing or order fulfilment.

### Example: Loan Processing Workflow

#### 1. Defined Tasks:


- **Task 1:** Verify Applicant Information
- **Task 2:** Perform Credit Check
- **Task 3:** Approval
- **Task 4:** Disbursement

#### 2. Task Dependencies:

- Task 2 (Credit Check) can only commence after the successful completion of Task 1 (Verify Applicant Information).
- Task 3 (Approval) is dependent on the positive outcome of Task 2 (Credit Check).

#### 3. Failure-Atomicity Requirements:

- Ensure that the workflow terminates in an acceptable state, even in the face of system failures.
- **For example:** If the application is approved, it undergoes a commitment process; if it fails to meet criteria, an abort mechanism is activated.



#### 4. Workflow Execution Components:

- **Scheduler:** Determines the task execution order based on predefined dependencies.
- **Recovery Routines:** Activated in case of failures, ensuring the restoration of the workflow state.

Transactional workflows are crucial for ensuring the **consistent and accurate completion of complex processes** like loan approval. They provide a framework where **tasks are logically sequenced, dependencies are managed, and failure scenarios are addressed**, guaranteeing the reliability and integrity of the overall workflow.

### 3. WHAT ARE REAL TIME TRANSACTION SYSTEM? EXPLAIN.

A real-time transaction system is a type of database system that is **designed to handle transactions within a strict timeframe dictated by the requirements of the system it supports**. These systems are used in environments where transactions must be processed within specific time to maintain the overall **functionality and correctness** of the system.

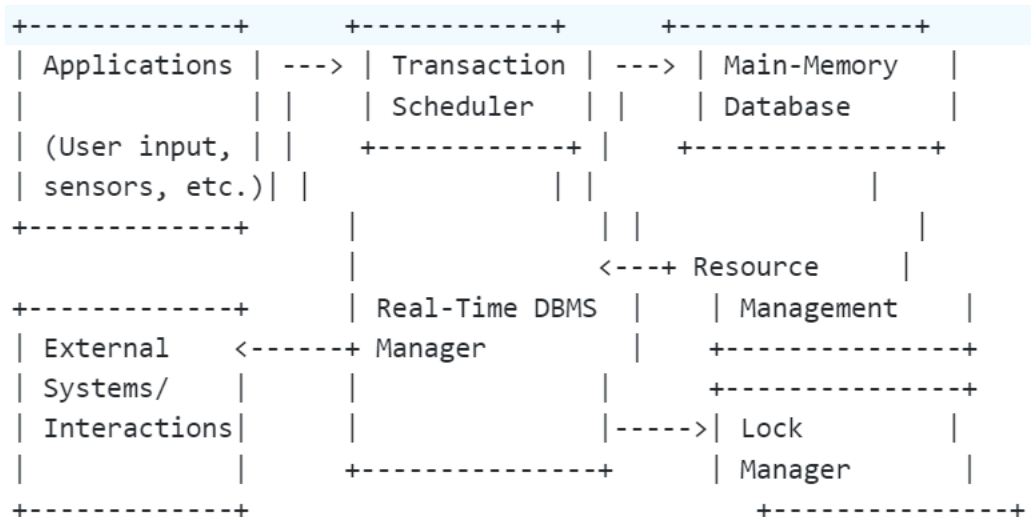
In real-time transaction systems, correctness is determined by **database consistency and the satisfaction of deadlines**. There are generally three types of deadlines considered in such a system:

1. **Hard Deadlines:** The **completion of a task within its deadline is crucial**. Failing to do so can lead to **catastrophic outcomes**. For example, in an air traffic control system, missing a hard deadline could result in significant safety risks.
2. **Firm Deadlines:** **Tasks that must be completed by their deadline to be of any value**. If a **task finishes after its firm deadline, it is worthless**. An example might be **online examinations**, where submissions after the deadline are not accepted.
3. **Soft Deadlines:** **These deadlines allow some flexibility**; For instance, **streaming live cricket match events** should ideally be as close to real-time as possible, **but short delays might still be acceptable**.

The challenges faced by real-time transaction systems include the **inherent variability of disk I/O operations, which can be unpredictable** and therefore complicate the management of transactions. To mitigate this, such systems often rely on **main-memory databases**, which store data in RAM instead of on slower disk-based storage, to facilitate faster access and processing times.

The architecture of a real-time transaction system **needs to be designed to ensure that it has adequate processing power to meet all deadlines**, without requiring excessive hardware resources which may not be economically feasible.

**Here's a high-level architecture of real-time transaction system:**



### The components in this architecture include:

- **Applications:** These **generate transactions** which might be initiated by user commands, sensor inputs, or interactions from external systems.
- **Transaction Scheduler:** Determines the order in which transactions are processed, taking into account the **priority and deadlines** associated with each transaction.
- **Real-Time DBMS Manager:** Oversees the general operation of the database system, ensuring that it **maintains consistency and meets the real-time requirements**.
- **Main-Memory Database:** Stores data in main memory for quick access, which is **essential for meeting strict time constraints**.
- **Resource Management:** **Handles the allocation of system resources such as CPU and memory** to ensure that transactions have the required resources to execute before their deadlines.
- **Lock Manager:** **Manages access to data items to prevent conflicts** and ensure **data consistency**, while also working within time constraints to prevent deadline violations.

In designing and operating a real-time transaction system, careful planning and management are critical to **ensuring that the system can consistently meet the time constraints it faces**, thereby providing reliable and timely service as expected by the system requirements.

### 4. GIVE THE DETAILS OF LONG DURATION TRANSACTIONS.

Long-duration transactions present unique challenges that differentiate them from the typical, shorter database transactions. They are a type of transaction where **user interaction and complex tasks with multiple steps form core components of the database operation**, often over a prolonged period. Below are the key details and challenges associated with long-duration transactions.

1. **Long Duration:** Such **transactions typically involve activities that cannot be completed quickly**. For example, an engineer working on a complex CAD model might have a design session that lasts for several hours or days. Traditional **concurrency control mechanisms**



aren't ideal for these scenarios because they generally assume transactions are short and can be completed quickly.

2. **Exposure of Uncommitted Data:** During long-duration transactions, it is often necessary for uncommitted data to be visible.
3. **Subtasks and Partial Rollback:** These transactions may have several intermediate states or checkpoints that users may wish to revert to without abandoning the entire transaction. In order to achieve this support for partial rollbacks is necessary.
4. **Recoverability:** This means that long-duration transactions require mechanisms that can recover uncommitted data in the event of a system crash, ensuring that users do not lose their progress.
5. **Performance:** Fast response times are crucial in systems that handle long-duration transactions. Users interact with these systems in real-time, and any lag can lead to a frustrating user experience and wasted time.

To manage these challenges, long-duration transactions are often modelled as nested transactions

- **Nested Transactions:** In this framework, the overall long-duration transaction is broken down into smaller transactions, some of which can be standard atomic database operations. These sub transactions can be committed independently, which allows for flexible management of data and resources without holding locks for extended time periods.
- **Transaction Failure Handling:** If a short-duration sub transaction within a long-duration transaction fails, it can be aborted without affecting the entire long-duration transaction. The ongoing longer transaction may be paused and resumed once the issues are resolved in the short-duration transactions.
- Apart from these there are other concurrency control protocols are used to resolve wait and abort problem with long duration transactions.
  1. Multi-version Protocol
  2. Optimistic Protocol

## 5. TRANSACTION MANAGEMENT IN MULTI-DATABASES

Transactional management in multi-database systems (MDBS) is complex due to the integration of multiple, autonomous local databases that may not be designed to work in conjunction. Their autonomy complicates the enforcement of a uniform concurrency control strategy across the entire system.

**Let's break down the key points:**

1. **Autonomy:** Each database in an MDBS is an independent entity with its own transaction management system. This autonomy principle means that local DBMSs do not directly

**communicate to coordinate transactions**, making global transaction management challenging.

2. **Local and Global Transactions:** There are two levels of transactions:

- **Local Transactions:** Managed by each local DBMS without MDBS interference.
- **Global Transactions:** Span across multiple databases and are managed by the MDBS.

3. **Concurrency Control Challenges:** To ensure the correctness of global transactions:

- **Global Two-Phase Locking (Global 2PL):** This **involves acquiring all necessary locks before any operation and only releasing them when the global transaction is fully complete**, ensuring serializability across databases.
- **Autonomy Limitations:** Since each site may have its own concurrency control, implementing a common strategy like strict 2PL across all local databases is not feasible.

4. **Solutions to Autonomy:**

- **Limiting Concurrent Execution:** Reducing levels of concurrency to a minimum can simplify transaction management but at the cost of performance.
- **Weaker Consistency Levels:** Allowing **less strict consistency models can permit more concurrency** but risks anomalies that serializability would usually prevent.

5. **Two-Level Serializability (2LSR):**

- **Local Serializability:** Each DBMS maintains its local serializability for local transactions.
- **Multidata base Serializability:** The MDBS strives for serializability among global transactions, **ignoring the impact of local transactions**.

6. **Read Protocols:**

- **Global-Read Protocol:** Global transactions can only **read local databases without update permissions**.
- **Local-Read Protocol:** Local transactions can read global data, but global transactions can't access local data.